

Aula 7

March 27, 2025

0.1 KDD

Base de dados -> descobrir padrões em dados e esse padrões ajudarem na tomada de decisão -> Descobrir conhecimento em dados (KDD)(insights)

Etapas do processo de KDD: - Dados estruturados - Pré-processamento - EDA - Análise Exploratória dos Dados - Aplica as técnicas e modelos e - Obtém insights ou gráficos

Etapas 1: Importando as bibliotecas

```
[2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy
```

```
[3]: # Carregando a base de dados
df = pd.read_csv('diabetes.csv')
```

```
[4]: print(df.head(10))
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0

6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                             768 non-null    int64
2   BloodPressure                       768 non-null    int64
3   SkinThickness                      768 non-null    int64
4   Insulin                            768 non-null    int64
5   BMI                                768 non-null    float64
6   DiabetesPedigreeFunction            768 non-null    float64
7   Age                                768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[6]: df.isnull().sum()
```

```
[6]: Pregnancies      0
      Glucose         0
      BloodPressure   0
      SkinThickness   0
      Insulin         0
      BMI            0
      DiabetesPedigreeFunction  0
      Age            0
      Outcome        0
      dtype: int64
```

```
[8]: df.isnull()
```

```
[8]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
..	
763	False	False	False	False	False	False	
764	False	False	False	False	False	False	

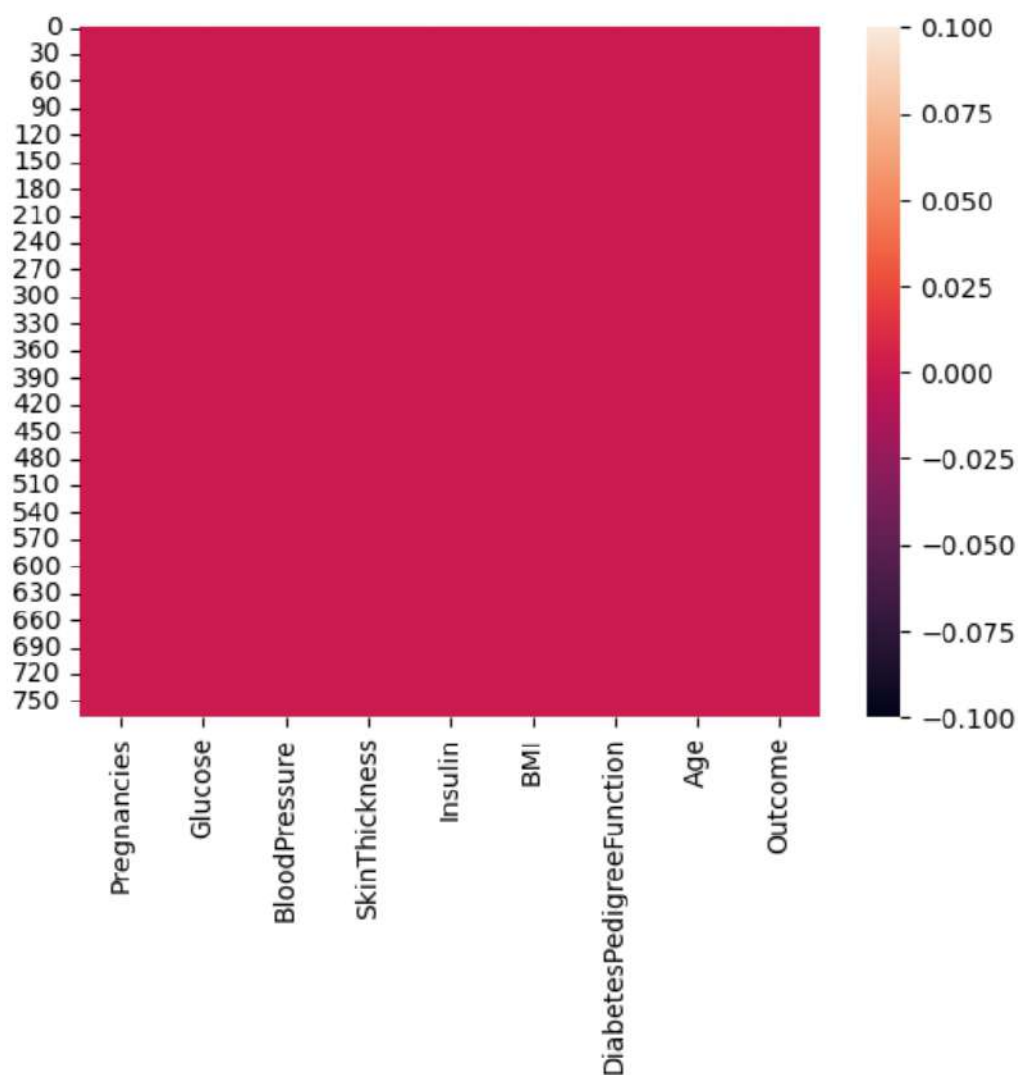
765	False	False	False	False	False	False	False
766	False	False	False	False	False	False	False
767	False	False	False	False	False	False	False

	DiabetesPedigreeFunction	Age	Outcome
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
..	""	""	""
763	False	False	False
764	False	False	False
765	False	False	False
766	False	False	False
767	False	False	False

[768 rows x 9 columns]

```
[7]: sns.heatmap(df.isnull())
```

```
[7]: <AxesSubplot: >
```



```
[9]: df.nunique()
```

```
[9]: Pregnancies      17
      Glucose         136
      BloodPressure    47
      SkinThickness    51
      Insulin         186
      BMI            248
      DiabetesPedigreeFunction  517
      Age            52
      Outcome         2
      dtype: int64
```

```
[11]: df['Pregnancies'].unique()
```

```
[11]: array([ 6,  1,  8,  0,  5,  3, 10,  2,  4,  7,  9, 11, 13, 15, 17, 12, 14])
```

```
[12]: df.describe().T
```


```
[12]:
```

	count	mean	std	min	25% \
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000
Glucose	768.0	120.894531	31.972618	0.000	99.00000
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000
Insulin	768.0	79.799479	115.244002	0.000	0.00000
BMI	768.0	31.992578	7.884160	0.000	27.30000
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375
Age	768.0	33.240885	11.760232	21.000	24.00000
Outcome	768.0	0.348958	0.476951	0.000	0.00000

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
[30]: sns.boxplot(df['SkinThickness'])
```

```
[30]: <AxesSubplot: >
```



```
[17]: Q1, Q3 = np.percentile(df['Pregnancies'], [25, 75])
```

```
[18]: Q1
```

```
[18]: 1.0
```

```
[19]: Q3
```

```
[19]: 6.0
```

```
[21]: # INTERVALO INTERQUARTIL
      IRQ = Q3-Q1
```

```
[23]: limite_inferior = Q1-(1.5*IRQ)
      limite_superior = Q3+(1.5*IRQ)
```

```
[24]: limite_inferior
```

```
[24]: -6.5
```

```
[25]: limite_superior
```

```
[25]: 13.5
```

```
[26]: df_limpo = df[(df['Pregnancies']>=limite_inferior)&(df['Pregnancies']<=limite_superior)]
```

```
[27]: sns
```

```
[27]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[764 rows x 9 columns]

```
[28]: sns.boxplot(df_limpo['Pregnancies'])
```

```
[28]: <AxesSubplot: >
```



```
[31]: Q1, Q3 = np.percentile(df['Insulin'], [25, 75])
      IRQ = Q3-Q1
      limite_inferior = Q1-(1.5*IRQ)
      limite_superior = Q3+(1.5*IRQ)
      df_limpo = □
      □df_limpo[(df_limpo['Insulin']>=limite_inferior)&(df_limpo['Insulin']<=limite_superior)]

      Q1, Q3 = np.percentile(df['Age'], [25, 75])
      IRQ = Q3-Q1
      limite_inferior = Q1-(1.5*IRQ)
      limite_superior = Q3+(1.5*IRQ)

      df_limpo = □
      □df_limpo[(df_limpo['Age']>=limite_inferior)&(df_limpo['Age']<=limite_superior)]

      Q1, Q3 = np.percentile(df['Glucose'], [25, 75])
      IRQ = Q3-Q1
      limite_inferior = Q1-(1.5*IRQ)
      limite_superior = Q3+(1.5*IRQ)
      df_limpo = □
      □df_limpo[(df_limpo['Glucose']>=limite_inferior)&(df_limpo['Glucose']<=limite_superior)]
```



```

Q1, Q3 = np.percentile(df['BloodPressure'], [25, 75])
IRQ = Q3-Q1
limite_inferior = Q1-(1.5*IRQ)
limite_superior = Q3+(1.5*IRQ)
df_limpo =
    df_limpo[(df_limpo['BloodPressure']>=limite_inferior)&(df_limpo['BloodPressure']<=limite_su

Q1, Q3 = np.percentile(df['SkinThickness'], [25, 75])
IRQ = Q3-Q1
limite_inferior = Q1-(1.5*IRQ)
limite_superior = Q3+(1.5*IRQ)
df_limpo =
    df_limpo[(df_limpo['SkinThickness']>=limite_inferior)&(df_limpo['SkinThickness']<=limite_su

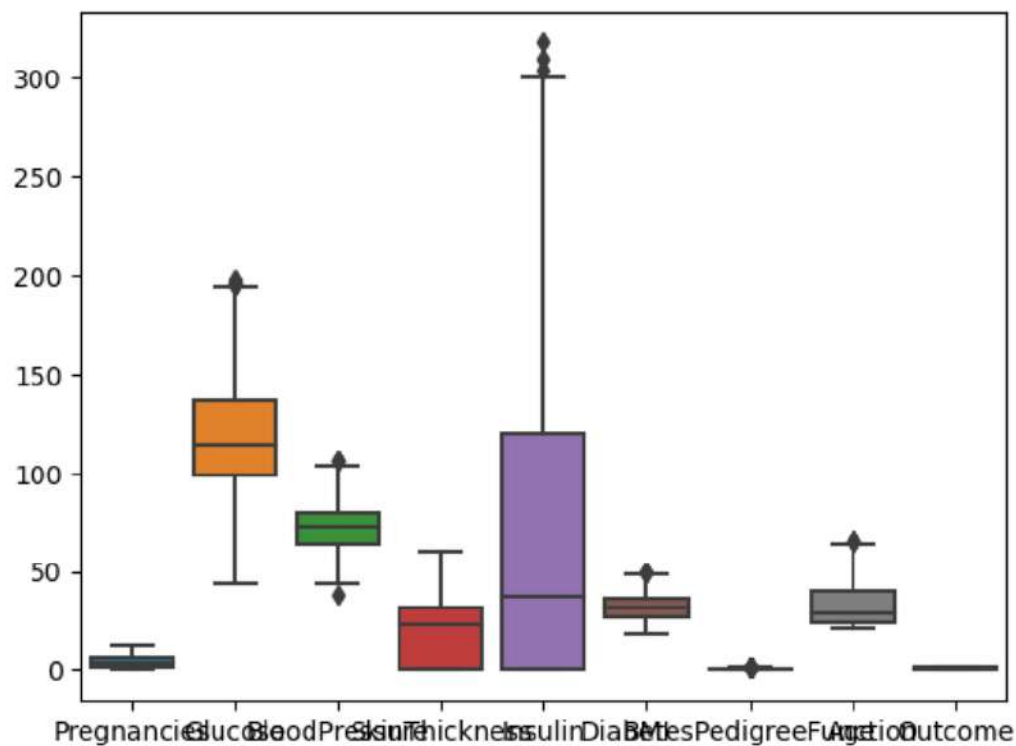
Q1, Q3 = np.percentile(df['BMI'], [25, 75])
IRQ = Q3-Q1
limite_inferior = Q1-(1.5*IRQ)
limite_superior = Q3+(1.5*IRQ)
df_limpo =
    df_limpo[(df_limpo['BMI']>=limite_inferior)&(df_limpo['BMI']<=limite_superior)]

Q1, Q3 = np.percentile(df['DiabetesPedigreeFunction'], [25, 75])
IRQ = Q3-Q1
limite_inferior = Q1-(1.5*IRQ)
limite_superior = Q3+(1.5*IRQ)
df_limpo =
    df_limpo[(df_limpo['DiabetesPedigreeFunction']>=limite_inferior)&(df_limpo['DiabetesPedigre

```

```
[32]: sns.boxplot(df_limpo)
```

```
[32]: <AxesSubplot: >
```



```
[33]: df_limpo
```

```
[33]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
5	5	116	74	0	0	25.6
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
5	0.201	30	0
..

763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[639 rows x 9 columns]

```
[34]: correlação = df.corr()
```

```
[35]: correlação
```

```
[35]:
```

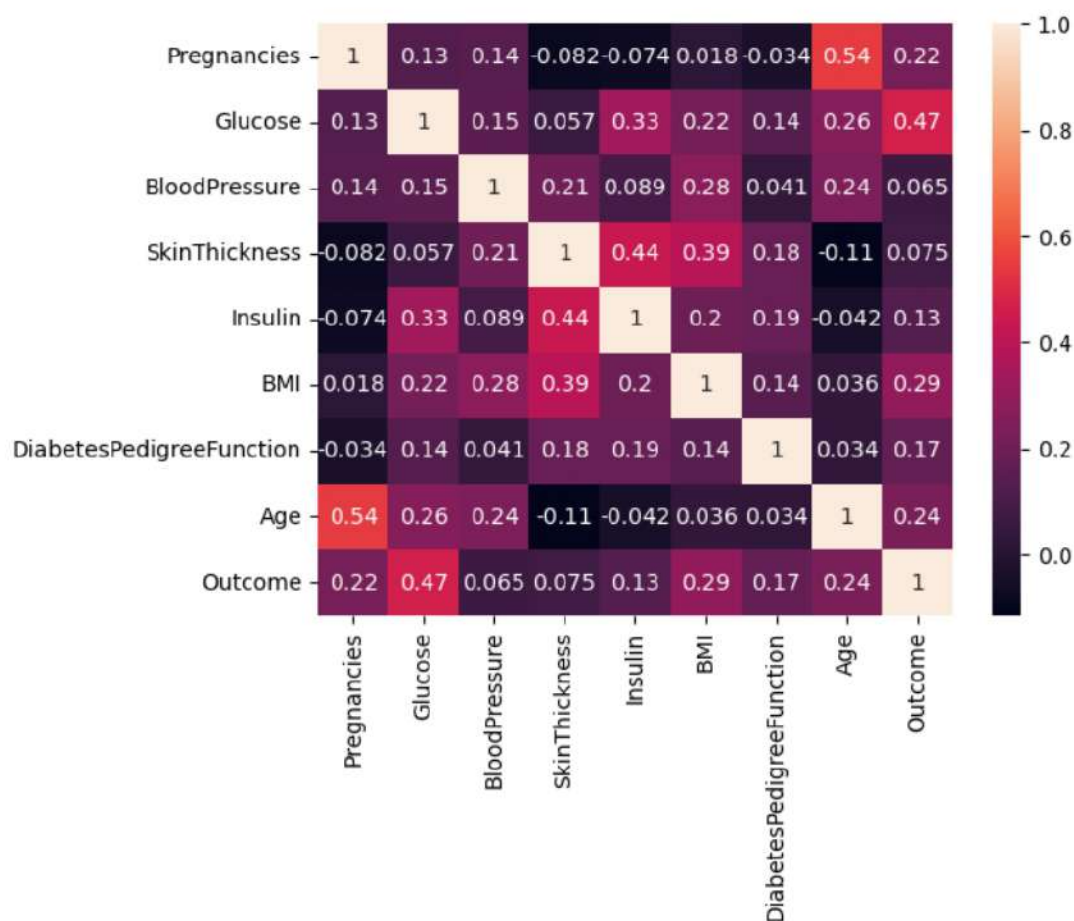
	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[38]: sns.heatmap(df.corr(), annot=True)
```

[38]: <AxesSubplot: >




```
[41]: correlação['Age'].sort_values(ascending=False)
```

```
[41]: Age                1.000000
Pregnancies            0.544341
Glucose                0.263514
BloodPressure          0.239528
Outcome                0.238356
BMI                   0.036242
DiabetesPedigreeFunction 0.033561
Insulin               -0.042163
SkinThickness         -0.113970
Name: Age, dtype: float64
```

```
[42]: plt.pie(df.Outcome.value_counts())
```

```
[42]: ([<matplotlib.patches.Wedge at 0x7fc9e8062bd0>,
      <matplotlib.patches.Wedge at 0x7fc9e804eb50>],
      [Text(-0.5025943242672991, 0.9784676515931925, ''),
      Text(0.5025944158780503, -0.9784676045369114, '')])
```



```
[45]: sns.pairplot(df, hue='Outcome')
```

```
[45]: <seaborn.axisgrid.PairGrid at 0x7fc9e076df50>
```




```
[46]: from sklearn.preprocessing import MinMaxScaler
```

```
[47]: escala = MinMaxScaler()
```

```
[48]: # X armazeno as variáveis independentes e Y a variável alvo que é a variável
      ↳ dependente
X = df.drop(columns=['Outcome'])
Y=df.Outcome
```

```
[49]: escala.fit(X)
Xescalar = escala.transform(X)
```

```
[51]: Xescalar
```

```
[51]: array([[0.35294118, 0.74371859, 0.59016393, ..., 0.50074516, 0.23441503,
          0.48333333],
          [0.05882353, 0.42713568, 0.54098361, ..., 0.39642325, 0.11656704,
          0.16666667],
          [0.47058824, 0.91959799, 0.52459016, ..., 0.34724292, 0.25362938,
          0.18333333],
          ...,
          [0.29411765, 0.6080402 , 0.59016393, ..., 0.390462 , 0.07130658,
          0.15      ],
          [0.05882353, 0.63316583, 0.49180328, ..., 0.4485842 , 0.11571307,
          0.43333333],
          [0.05882353, 0.46733668, 0.57377049, ..., 0.45305514, 0.10119556,
          0.03333333]])
```

```
[52]: df.describe().T
```

```
[52]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
[53]: df_titanic=pd.read_csv('titanic.csv')
```

```
[54]: df_titanic.head()
```

```
[54]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	\
0	0	3	male	22.0	1	0	7.2500	S	Third	
1	1	1	female	38.0	1	0	71.2833	C	First	
2	1	3	female	26.0	0	0	7.9250	S	Third	
3	1	1	female	35.0	1	0	53.1000	S	First	

```
4          0          3    male  35.0          0          0  8.0500          S  Third
```

```
      who  adult_male deck  embark_town alive  alone
0    man          True  NaN  Southampton    no  False
1  woman          False   C    Cherbourg   yes  False
2  woman          False  NaN  Southampton   yes   True
3  woman          False   C    Southampton   yes  False
4    man          True  NaN  Southampton    no   True
```

```
[57]: df_titanic.duplicated()
```

```
[57]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
      886     True
      887     False
      888     False
      889     False
      890     False
      Length: 891, dtype: bool
```

```
[58]: df_titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   survived        891 non-null    int64
1   pclass          891 non-null    int64
2   sex             891 non-null    object
3   age            714 non-null    float64
4   sibsp          891 non-null    int64
5   parch          891 non-null    int64
6   fare           891 non-null    float64
7   embarked       889 non-null    object
8   class          891 non-null    object
9   who            891 non-null    object
10  adult_male     891 non-null    bool
11  deck           203 non-null    object
12  embark_town    889 non-null    object
13  alive          891 non-null    object
14  alone          891 non-null    bool
dtypes: bool(2), float64(2), int64(4), object(7)
memory usage: 92.4+ KB
```



```
[59]: cat_col = [col for col in df_titanic.columns if df_titanic[col].dtype ==  
↳ 'object']  
print(cat_col)
```

```
['sex', 'embarked', 'class', 'who', 'deck', 'embark_town', 'alive']
```

```
[60]: num_col = [col for col in df_titanic.columns if df_titanic[col].dtype !=  
↳ 'object']  
print(num_col)
```

```
['survived', 'pclass', 'age', 'sibsp', 'parch', 'fare', 'adult_male', 'alone']
```

```
[61]: df_titanic[cat_col].nunique()
```

```
[61]: sex                2  
embarked              3  
class                3  
who                  3  
deck                 7  
embark_town          3  
alive                2  
dtype: int64
```

```
[62]: df_titanic['deck'].unique()[:50]
```

```
[62]: array([nan, 'C', 'E', 'G', 'D', 'A', 'B', 'F'], dtype=object)
```

```
[63]: df_titanic.isnull().sum()
```

```
[63]: survived          0  
pclass                0  
sex                  0  
age                  177  
sibsp                 0  
parch                 0  
fare                 0  
embarked              2  
class                 0  
who                   0  
adult_male            0  
deck                 688  
embark_town           2  
alive                 0  
alone                 0  
dtype: int64
```

```
[64]: 177/890
```

```
[64]: 0.19887640449438201
```

```
[65]: 2/890
```

```
[65]: 0.0022471910112359553
```


```
[66]: df_titanic_2=df_titanic.drop(columns='deck')  
df_titanic_2.dropna(subset=['embark_town'],axis=0,inplace=True)
```

```
[67]: df_titanic_2.shape
```

```
[67]: (889, 14)
```

```
[73]: sns.boxplot(data=df_titanic_2, x='pclass', y='age')
```

```
[73]: <AxesSubplot: xlabel='pclass', ylabel='age'>
```



0.2 Exercício

Substitua os valores faltantes do atributo age da base de dados pela mediana presentes nos dados da primeira, segunda e terceira classes da base de dados titanic, conforme explicado durante a aula.